ENGINEERING AUDIT · NOVEMBER 2026

SaaS Shield Complete Product and Production Readiness Audit

Competitor intelligence, AI analysis and evidence monitoring built on Node.js 22, Express and SQLite/PostgreSQL. This report states what was measured, what was mocked and what has never run.

Field	Value
Repository	bittubuilds/saas-shield
Branch / HEAD	feat/db-durability @ cf53813 (verified before audit)
Working tree	10 modified source and test files, 1 untracked E2E test, uncommitted
Production impact	None. Nothing committed, pushed, merged or deployed. Render and Neon untouched.
Live AI call	Not made. Provider path verified as code, not as behaviour.
Test evidence	SQLite 273 tests / 269 pass / 0 fail / 4 skipped; PostgreSQL integration 5/5; durability 16/16; whole suite on PostgreSQL 262 pass / 10 fail (all diagnosed)

Verdict in one paragraph

The pipeline is real: crawling, deterministic extraction, evidence versioning, incremental AI analysis, synthesis, versioned reports and PDF export all execute and are covered by tests. The AI transport is a genuine HTTP client for Groq with timeouts, retries, quota pausing, citation enforcement and duplicate-call prevention, but no completion has ever been sent to the real provider, so every AI behaviour is mock-verified only. The decisive product gap is surface: the frontend bundle contains no call to `/ai/*` or `/competitors/*`, so competitor intelligence and AI findings exist only as API endpoints. Phase 1 is complete as an engine and incomplete as a product.

SECTION 1

Executive summary

Three findings matter more than the rest, and all three are about verification rather than volume of code.

#	Finding	Evidence and explanation	Status
1	The AI provider path has never executed.	Groq is genuinely integrated: `callGroq()` issues a real fetch to `api.groq.com/openai/v1/chat/completions` with bearer auth. A key is present in this environment and a keyless probe returned HTTP 401, proving the endpoint is live. No completion, no JSON-mode response and no token accounting has ever been produced by the real model.	UNVERIFIED
2	Half the product has no user interface.	The shipped bundle calls auth, workspaces, websites, scans, results, changes, monitoring, alerts, discoveries, deterministic export and settings. It never calls the AI, competitor or screenshot APIs. A user cannot see a finding, a competitor profile, a screenshot or an AI report.	FAIL
3	Everything else holds up.	Isolation, quota pausing, idempotency, versioning, SSRF, injection fencing and both database backends are implemented and tested. The 10 whole-suite PostgreSQL failures are test-bank design, not product defects.	PASS

What would change this verdict: one live provider completion, and a frontend that renders the findings the API already stores. Neither is a rewrite.

SECTION 2

Real AI integration verification

Four states are kept separate throughout this report: **code implemented**, **mock-tested**, **real provider integration tested**, **production configuration verified**. Only the first two are satisfied for AI behaviour.

Check	Finding	Evidence	Status
Provider implemented	Groq, OpenAI-compatible chat completions API	src/ai/provider.js:78-110	PASS
Genuine or mock	Genuine `fetch` with bearer auth. `mock` exists separately and never touches the network	src/ai/provider.js (callGroq, callMock)	PASS
Model configured	`openai/gpt-oss-120b`, base `https://api.groq.com/openai/v1`	src/ai/provider.js:26-27	PASS
Adapter location	Single file owns the wire format; callers speak {system,user} to parsed JSON	src/ai/provider.js:1-16, complete() :137	PASS
Credential loading	`AI_API_KEY` or `GROQ_API_KEY` from env, server-side only	src/ai/provider.js:28	PASS
Credential validation	`isConfigured()` gates routes with 503; key never returned or logged; errors redacted	provider.js:33-36, 44-48; tests/ai.test.js:148	PASS
Real HTTP requests	Code yes. Endpoint reachable: keyless GET /models returned 401	curl probe, zero quota spent	PARTIAL
Response parsing	Strict JSON: fences stripped, object required, else `invalid_json` hard failure	provider.js:62-76; tests/ai.test.js:131	PASS
Timeouts	45s AbortController per request, abort mapped to retryable timeout	provider.js:112-142	PASS
Retries	2 retries, exponential backoff capped at 8s, transient only	provider.js:137-152	PASS
Rate limit / failure mapping	429 quota_exhausted (not retried), 401/403 auth_failed, 404 model_unavailable, 5xx provider_outage	provider.js:54-61	PASS
Quota exhaustion pauses safely	Stage paused, attempt budget refunded, completed stages kept, scan reported partial	engine.js:243-262; tests/ai.test.js:425	PASS

Check	Finding	Evidence	Status
Usage limits enforced	40 requests/scan, 200/day/workspace, 6 competitors, 3 pages, 3 attempts	engine.js:22-31; engine.js:196-207	PASS
Duplicate calls prevented	Atomic claim before the model call; unique index makes the loser adopt the winner; live runs respected	engine.js:209-223; store.js:51; tests/ai.test.js:365	PASS
Incremental: unchanged evidence costs nothing	Input hash over evidence ids, content hashes, analysis and prompt versions, plus structured fingerprint	engine.js:56-64; tests/ai.test.js:205	PASS
Synthesis uses structured evidence	Reads stored findings only, never pages	src/ai/synthesize.js:44-50	PASS
Findings cite valid evidence	Uncited findings dropped; forged ids dropped; foreign-workspace ids filtered	schema.js:34-50; engine.js:100-118; tests/ai.test.js:154	PASS
Historical reports preserved	`reportByVersion`, `reportHistory`, `?version=N` on JSON and PDF	store.js:198-213; ai-routes.js:103-160	PASS
Live completion executed	Never. Key present, no completion sent, no response shape observed	no call made; see Section 12	UNVERIFIED

Reference limits and cost model

Documented free-tier ceiling for this model is 30 requests per minute, 1,000 per day, 8,000 tokens per minute and 200,000 per day. The code is shaped around the 8,000 tokens-per-minute wall: one page of roughly 4,000 characters per request, 1,400 output tokens cap. A full scan runs 4 stages per competitor over up to 3 pages, so 12 to 24 requests plus one synthesis, which against the 200/day workspace ceiling allows roughly eight scans per workspace per day.

Known determinism caveat

`temperature` is 0.1 with no seed, so identical evidence can yield different findings. Cost idempotency is unaffected (the hash is over evidence, not output), but output is not reproducible and no test pins it.

SECTION 3

AI output quality assessment

Every quality claim below rests on deterministic mock responses. There is no golden record captured from the real model, so these results prove the validation boundary, not the model.

Dimension	What was verified	Evidence	Status
Competitor positioning	positioning and differentiator findings, brief-driven	prompts.js:42; engine.js:141-150	PASS
Feature comparison	common_features and unique_features in synthesis	schema.js:97-118	PASS
Pricing comparison	a plan with no stated price stays literal "unknown"	schema.js:56; tests/ai.test.js:162	PASS
Advantages and weaknesses	SWOT buckets whitelisted; invalid bucket dropped	schema.js:118-124; tests/ai.test.js:168	PASS
Market gaps	gaps[] carries evidence, confidence and wording	prompts.js:135; schema.js:110-115	PASS
Actionable recommendations	grounded in findings but the link is unchecked prose	prompts.js:130-137; schema.js:126-134	PARTIAL
Evidence citations	every finding must cite a supplied id or it is discarded	schema.js:34-50	PASS
Confidence and uncertainty	values clamped 0-1, claim level whitelisted, no calibration	schema.js:22-28	PARTIAL

Dimension	What was verified	Evidence	Status
Hallucination prevention	invented competitors dropped, absent facts become unknown, uncited claims refused	schema.js:84-96	PASS
Unknown prices remain unknown	price falls back to the string "unknown", never derived	schema.js:58	PASS
Unsupported features not asserted	claim level distinguishes explicit, inferred and unverified	prompts.js:28; schema.js:41	PASS
Similarity not described as copying	system rule 6 and synthesis rule 5 forbid plagiarism assertions	prompts.js:30, :127	PASS
Recommendations grounded	input is findings only, so it cannot exceed them; wording is free	synthesize.js:44-50	PARTIAL
Injection resistance	random-UUID fence, overriding system rule, citable-id whitelist; model can still be nudged in prose but cannot write unsupported fact	prompts.js:12-31, :103-120; tests/ai.test.js:273-282	PASS

SECTION 4

Complete product functionality checklist

The full user workflow, checked against the code and against the 15-step end-to-end test that runs the pipeline with deterministic fixtures. "API" means the capability exists and is tested at the route or service layer but is not reachable from the shipped interface.

#	Capability	Finding	Evidence	Status
1	Registration and login	Scrypt hashing, opaque session tokens stored as SHA-256, httpOnly SameSite cookies	src/routes/auth-routes.js; src/util.js:14; src/auth.js:9-25	PASS
2	Add a website	Route with 30/min limiter, normalisation, workspace scope	website-routes.js:47; E2E step 1	PASS
3	Discover competitors	Search-backed discovery via injected search provider	src/discover.js	PARTIAL
4	Identify and score competitors	Taxonomy classification plus weighted scoring with threshold	src/competitor/score.js; E2E step 2	PASS
5	Reject irrelevant candidates with reasons	Rejection reason persisted for audit	src/competitor/identify.js; migrations 8	PASS
6	Crawl competitor pages	Deep crawl, ranked discovery, bounded depth, same-site enforcement	src/crawler.js; src/competitor/deep-crawl.js	PASS
7	Classify pages	Path and content classification	src/competitor/classify.js; E2E steps 3-6	PASS
8	Extract structured features	Deterministic claim levels, no model involved	extract-features.js	PASS
9	Extract structured pricing	Original wording kept, unknown as NULL, no currency conversion	extract-pricing.js	PASS
10	Store evidence and versions	Content hashing ignores whitespace and case; version chains per identity	evidence-versions.js; E2E steps 7-8	PASS
11	Track freshness	Five-state freshness on 24h and 168h TTLs	src/competitor/freshness.js; E2E step 9	PASS
12	Capture screenshots	Real Playwright capture, bounded pool of 3, 30s timeout, dedupe by hash. Unavailable in this host: no playwright package installed	src/screenshots/browser.js; src/ai/vision.js	PARTIAL
13	Incremental AI analysis	Stage-level claim, resume, skip on unchanged evidence	src/ai/engine.js; E2E steps 11-12	PASS

#	Capability	Finding	Evidence	Status
14	Cross-competitor synthesis	Hash-claimed report, one model call, explicit coverage gaps	src/ai/synthesize.js; E2E steps 13-15	PASS
15	Reports and PDF export	Deterministic report PDF plus AI report PDF, both versioned	src/pdf/ai-report-pdf.js; src/pdf/engine.js	PASS
16	Access historical reports	JSON and PDF by explicit version	ai-routes.js:103-160; store.js:198	PASS
17	Recover interrupted jobs	Stale analysis claims, stale report claims, scan recovery on boot	engine.js:209-223; src/scan-recovery.js	PASS
18	Enforce isolation and usage limits	Workspace scoping at query level, 404 across tenants, request ceilings	ai-routes.js; tests/ai.test.js:332; E2E isolation step	PASS

Interface gap, stated precisely: the shipped bundle contains 152 escaped DOM writes and calls 33 API paths, none of which is an AI, competitor, screenshot or findings endpoint. Feature rows 4, 5, 8, 9, 12, 13, 14, 15 and 16 are therefore backend-only. For a user, that half of the product does not exist yet.

SECTION 5

Security audit findings by severity

#	Severity	Finding	Detail	Evidence	Impact
1	MEDIUM	No per-user AI quota	The daily ceiling is per workspace.	src/ai/store.js:154	Members of one workspace share 200 requests/day with no per-person limit.
2	MEDIUM	AI analysis is not triggered automatically	Only an explicit POST starts a run.	ai-routes.js:64; no caller elsewhere in src	Status payload can report ai_configured while nothing analyses a finished scan.
3	MEDIUM	SSRF suite is configuration-fragile	Exporting SHIELD_ALLOW_PRIVATE_IPS=true turns 6 SSRF assertions into failures.	observed: 16 failures instead of 10	Protection is correct; the test bank assumes default config globally.
4	LOW	No CSRF token	Mitigated by SameSite cookie.	src/auth.js:203-210	Acceptable at current scope; would matter if SameSite is relaxed.
5	LOW	Stale security documentation	Docs claim no screenshots, no AI route tests, no AI PDF.	docs/AI_SYSTEM.md:222-235	Misleading for operators assessing exposure.
6	INFO	Error surfaces are sanitized	Key redacted by pattern and by literal value.	provider.js:44-48; tests/ai.test.js:148	No leak found.

Controls verified as present

Scrypt password hashing with per-user salt and parameterised cost; opaque session tokens with only their SHA-256 stored; httpOnly, SameSite and secure cookies; per-route rate limits on auth, scan, discovery, monitoring, collect, notification and AI analysis; DNS-revalidated SSRF checks including the 169.254.169.254 metadata address and IPv6 literals; manual redirect revalidation on every hop; timing-safe cron secret; Svix signature verification with 300-second replay tolerance for email webhooks; SQL identifiers from whitelists with parameterised values; escaped interpolation for all untrusted text in the interface. Screenshot image routes are workspace-scoped and resolve the record before touching storage.

SECTION 6

Database and reliability assessment

Area	Finding	Evidence	Status
SQLite compatibility	Full suite green	273 tests / 269 pass / 0 fail / 4 skipped	PASS

Area	Finding	Evidence	Status
PostgreSQL compatibility	Supported path green on PG 15	tests/postgres.integration.test.js 5/5	PASS
Migration consistency	13 migrations, applied list derived from the module	src/db/migrations.js; pg_advisory_xact_lock and BEGIN IMMEDIATE	PASS
Foreign keys and indexes	121 REFERENCES and CREATE INDEX statements; partial unique index on running reports	src/db/migrations.js:1-13	PASS
Concurrent job execution	Atomic claim wins exactly once; two-process migration and analysis tested	tests/ai.test.js:365; db-durability PG tests	PASS
Backup and restore	Snapshot verify, swap sidecars, parent-before-child restore, real PG savepoint	src/db/backup.js; durability 16/16	PASS
Restart recovery	Stale scan recovery on boot, stale claim reclaim window	src/scan-recovery.js; tests/scan-reliability.test.js	PASS
Connection failure behaviour	Startup validation, dial errors surfaced, no silent fallback	src/db/index.js; src/db/pg-backend.js	PASS
Partial scan recovery	Completed stages retained, scan reported partial	engine.js:268-282	PASS
Data retention and report history	Report versions retained; alert retention pruning	store.js:198-213; src/alerts.js retention	PASS
Whole-suite PostgreSQL run	10 failures, all diagnosed as test-bank design	see below	PARTIAL

Whole-suite PostgreSQL failures, diagnosed

The suite is designed around one isolated database per test file. When DATABASE_URL is set, every file shares one PostgreSQL database, which produces two distinct effects. None of them is a product defect, and the supported PostgreSQL path (integration plus durability with TEST_DATABASE_URL) is fully green.

Affected	Observation	Diagnosis
9 tests in db-durability and scan-reliability	Assert a SQLite file backend: snapshot verification, swap restore, deleted-row restore, cross-process reads, two-process migration, restart recovery	These tests name SQLite in their own titles; they cannot pass against a PostgreSQL backend by construction
1 test in monitor.test.js	Expects 0 due checks for its workspace, receives 4	Rows created by other test files are visible in the shared database: test isolation, not scheduler behaviour

Caveat: the monitor failure is strongly indicated as contamination but was not isolated by running that file alone on a fresh database, because the audit ran out of step budget at that point. Treat it as the least-certain diagnosis in this report.

SECTION 7

Performance and scalability findings

No benchmark is quoted here because none was reproducible in this environment. What follows is derived from configuration constants and call structure, which are facts in the code.

Dimension	What the code does	Realistic implication	Evidence
AI request volume per scan	4 stages x up to 6 competitors, 1 page per request, plus 1 synthesis	12 to 24 analysis requests plus 1	engine.js:22-31; prompts.js:57
Realistic daily capacity	200 requests/day workspace ceiling against 13 to 25 per scan	roughly 8 scans per workspace per day	engine.js:26
Prompt size ceiling	4,000 characters of page text per request, 1,400 output tokens	cross-page context is lost by construction	prompts.js:57; provider.js:96

Dimension	What the code does	Realistic implication	Evidence
Synthesis coverage	Findings trimmed to 12 per competitor before prompting	large scans are summarised, not compared exhaustively	synthesize.js:27
Crawler concurrency	Bounded deep crawl with depth and page caps	bounded by design, not measured	src/competitor/deep-crawl.js
Browser resource use	Pool of 3, 30-second timeouts, failures non-blocking	bounded; unavailable without playwright	src/screenshots/browser.js
Duplicate crawling	URL dedupe and content hashing before storage	same page not stored twice	evidence-versions.js
Query efficiency	Workspace-scoped indexed lookups; no unbounded table scans in the AI path	adequate at single-workspace scale	src/ai/store.js
Large report handling	PDF rendered in memory per request	a very large history grows response time	src/pdf/ai-report-pdf.js

SECTION 8

Phase 1 milestone-by-milestone completion

Milestones follow the Phase 1 build sequence; each is judged on executed behaviour, not on files existing. A milestone carrying a blocked or unverified capability cannot be marked complete.

Id	Milestone	Basis for the judgement	Status
M1	Audit, stabilisation, security review	Race conditions in the migration runner fixed, regression baseline established	COMPLETE
M2	Database durability and migrations	Persistent backend, backup and restore, idempotent migrations, both engines	COMPLETE
M3	Competitor intelligence foundation	Competitor, discovery, page and evidence tables; deterministic classify, identify, score	COMPLETE
M4	Collection and extraction	Deep crawl, page classification, feature and pricing extraction	COMPLETE
M5	AI intelligence engine	All four stages, schema validation, quota and retry handling	PARTIAL
M6	Evidence versioning and incrementality	Content hashing, version chains, freshness, unchanged evidence costs nothing	COMPLETE
M7	Synthesis, reports and export	Versioned synthesis, JSON history, AI PDF, deterministic PDF parity	COMPLETE
M8	Product integration	Frontend surface for intelligence, findings and reports	NOT IMPLEMENTED

M5 is PARTIAL, not complete, for one reason: the engine has never spoken to the provider it is built for. Every other M5 property (validation, idempotency, quota, resume) is proven.

Completion percentage, with the calculation shown

Component	Basis	Contribution
18 workflow capabilities in Section 4	15 full, 3 partial (discovery, screenshots, and steps that are API-only)	14.5 / 18 = 80.6%
Cross-cutting gate: live AI unverified	AI capability scored at 50% weight until a real completion succeeds	applied to M5 only
Cross-cutting gate: no intelligence interface	Rows 4, 5, 8, 9, 12, 13, 14, 15 and 16 are unreachable by a user	product surface judged separately
Reported figure	Engine completeness 81%, product completeness materially lower	Phase 1 = 78% (transparent, not a benchmark)

The 78% is an explicit engineering judgement, not a measurement. The defensible measured numbers are the test results in Section 6 and the 18-capability checklist in Section 4; the percentage condenses them and should be read that way.

SECTION 9

Remaining Phase 1 tasks

Task	What completion requires	Status	Reference
Complete M5	One controlled live provider completion, then fix whatever it exposes	UNVERIFIED	Section 2
Complete M8	Findings, competitor profile, screenshot viewer and AI report screens in the frontend	FAIL	Section 4
Close discovery	Configure and verify a real search provider so discovery is not fixture-only	PARTIAL	src/discover.js
Close screenshots	Install Playwright in the deployment image, then verify capture end to end	PARTIAL	src/screenshots/browser.js
Trigger analysis automatically	Run AI after scan completion, with the same claim and quota guards	PARTIAL	ai-routes.js
Per-user AI quota	Move the daily ceiling from workspace to user, or add a second gate	FAIL	src/ai/store.js:154
Refresh documentation	AI_SYSTEM.md still lists implemented features as limitations	FAIL	docs/AI_SYSTEM.md:222
PostgreSQL suite isolation	One database or schema per test file so the suite gates both backends	PARTIAL	Section 6

SECTION 10

Phase 2 readiness for market validation

Payments and paid subscriptions are explicitly out of scope until validation. These are the prerequisites for putting this in front of real founders.

Prerequisite	State	Status
Product onboarding	A new user can add a site and get a result, but cannot reach intelligence, findings or the AI report	PARTIAL
User experience	Eight-item navigation was already flagged as confusing; the richest output has no screen at all	PARTIAL
Competitor discovery accuracy	Deterministic scoring and rejection reasons exist; real search is unconfigured, so accuracy is unmeasured on live data	PARTIAL
Report usefulness	Sections, provenance and explicit unknowns are implemented; never evaluated by a human reader of a real report	PARTIAL
AI reliability	Guards are solid, model behaviour unverified, no live latency or failure data	UNVERIFIED
Usage limits	Per-scan and per-workspace ceilings implemented; per-user quota absent	PARTIAL
Error messages	Provider errors classified and sanitized; 409/503 states explicit for missing configuration	PASS
Documentation	Deployment, durability, screenshots and AI docs exist; AI doc is stale on three points	PARTIAL
Deployment configuration	Render blueprint present; Neon cutover checklist present; production connection still pending	PARTIAL
Monitoring and operational visibility	Health endpoint and internal status route expose backend and SSRF mode; no metrics, no error aggregation	PARTIAL
Backup strategy	Snapshot, verify, restore and retention implemented and tested on both engines	PASS

Prerequisite	State	Status
Privacy and data handling	Public content only; no per-user analytics; retention pruning for alerts; no published privacy note yet	PARTIAL
Feedback collection	Not implemented	FAIL

SECTION 11

Critical blockers and risks

#	Blocker	Why it matters	Severity	Smallest unblocking step
1	No live AI verification	Every AI claim in the product rests on mock behaviour. The first real call may surface schema drift, token ceiling problems or JSON-mode differences.	High	One request, negligible quota
2	No intelligence interface	Users cannot see the feature that differentiates the product. Backend work is invisible.	High	Frontend work, no new backend
3	Discovery depends on an unconfigured provider	Competitor discovery is fixture-only in practice.	Medium	Provider configuration plus a live check
4	Screenshots unavailable in deployment	Visual evidence is documented but not capturable here.	Medium	Add playwright to the image
5	Whole-suite PostgreSQL gate unusable	The suite cannot be a CI gate on both backends until files get isolated databases.	Medium	Per-file database isolation

SECTION 12

Prioritized recommended next steps

Order	Action	Rationale	Note
1	One controlled live AI completion	Converts the largest UNVERIFIED into evidence or a defect list. Approximately 10 to 30 tokens, one of 1,000 daily free-tier requests, no monetary cost.	Approval required before running
2	Frontend surface for findings and the AI report	Turns completed backend work into a product. Highest user-visible value per unit of work.	No dependency on step 1
3	Trigger analysis on scan completion	Removes a hidden requirement that a user know an endpoint exists.	Reuses existing claim and quota guards
4	Per-user AI quota	Prevents one workspace member exhausting the shared daily ceiling.	Small change in store.js
5	Per-file PostgreSQL test isolation	Makes the suite a real gate on both backends.	Test harness only
6	Playwright in the deployment image	Closes the screenshot capability.	Deployment change, needs authorisation
7	Documentation refresh	AI_SYSTEM.md currently understates the system in three places.	Low risk

Evidence base for this report

Item	Result
Repository state	git rev-parse HEAD = cf53813; working tree dirty with the 10 modified files listed in the cover
SQLite suite	npm test: 273 tests, 269 pass, 0 fail, 4 skipped (PostgreSQL-only tests)
PostgreSQL integration	node --test tests/postgres.integration.test.js with DATABASE_URL: 5/5 pass
PostgreSQL durability	node --test tests/db-durability.test.js with TEST_DATABASE_URL: 16/16 pass, including real savepoint, restore ordering and two-process migration

Item	Result
Whole suite on PostgreSQL	npm test with DATABASE_URL: 277 tests, 262 pass, 10 fail, 5 skipped, all 10 diagnosed in Section 6
Provider reachability	curl https://api.groq.com/openai/v1/models without credentials returned 401, proving the endpoint is live and requires auth. No quota consumed, no key transmitted.
Not measured	Live model output quality, latency, real search discovery accuracy, screenshot capture in this host, any throughput benchmark

Audit limitations

Limitation	Consequence
Live model behaviour	No completion was sent to Groq. Output quality, JSON-mode compliance, latency and token accounting are therefore unmeasured. Approval is required before that call.
Performance	No benchmark was run. Throughput figures in Section 7 are arithmetic over configuration constants, labelled as such.
One diagnosis is not isolated	The monitor test failure is attributed to shared-database contamination by construction, not by running that file alone on a fresh database.
Human evaluation	No reader has judged a real AI report for usefulness. Report quality is assessed structurally, not editorially.
Production state	Render and Neon were not inspected or modified. Production configuration is unverified beyond what the repository declares.
Files touched by this audit	This report and the script that generates it, both under reports/. No source, test or configuration file was modified, committed or pushed.

Generated 7 November 2026 from the repository at cf53813 on branch feat/db-durability. Audit only: no fix was implemented, nothing was committed or deployed, and no credential appears in this document.